

PDR for k -safety

No Author Given

No Institute Given

Abstract. We introduce an algorithm for k -safety. Based on PDR.

1 Introduction

2 Preliminaries

In this section, we present notations and background that is required for the description of our algorithm.

Safety verification. A transition system T is a tuple $(X, Init, Tr, Bad)$, where X is a set of variables that defines the states of the system (i.e., 2^X), $Init$ and Bad are formulas with variables in X denoting the set of initial states and bad states, respectively, and Tr is a formula with free variables in $X \cup X'$, denoting the transition relation. A state $s \in 2^X$ is said to be reachable in T if and only if (iff) there exists a state $s_0 \in Init$, and $(s_i, s_{i+1}) \in Tr$ for $0 \leq i \leq N$, and $s = s_N$. For simplicity of presentation, we assume that the initial states do not include bad states, that is $Init \Rightarrow \neg Bad$.

An *inductive invariant* is a formula Inv that satisfied:

$$Init(X) \Rightarrow Inv(X) \quad Inv(X) \wedge Tr(X, X') \Rightarrow Inv(X') \quad (1)$$

A transition system T is SAFE iff there exists an inductive invariant Inv s.t. $Inv(X) \Rightarrow \neg Bad(X)$. In this case we say that Inv as a *safe* inductive invariant. A formula Inv is said to be inductive *relative* to a formula F if Inv satisfies (1) with $Tr(X, X')$ replaced by $F(X) \wedge Tr(X, X')$.

A transition system T is UNSAFE iff there exists a state $s \in Bad$ s.t. s is reachable. Equivalently, T is UNSAFE iff there exists a number N such that the following formula is satisfiable:

$$Init(X_0) \wedge \left(\bigwedge_{i=0}^{N-1} Tr(X_i, X_{i+1}) \right) \wedge Bad(X_N) \quad (2)$$

where $X_i = \{a_i \mid a \in X\}$ is a copy of the variables used to represent the state of the system after the execution of i steps. We use $BMC(T, k)$ to denote a procedure that checks satisfiability of Formula 2. When T is UNSAFE and $s_N \in Bad$ is the reachable state, the path from $s_0 \in Init$ to s_N is called a *counterexample* (CEX).

The *safety verification* problem is to decide whether a transition system T is SAFE or UNSAFE, i.e., whether there exists a safe inductive invariant or a counterexample.

Input: A safety problem $\langle \text{Init}(X), \text{Tr}(X, X'), \text{Bad}(X) \rangle$.
Output: *Unreachable* or *Reachable*
Data: A cex queue \mathcal{Q} , where $c \in \mathcal{Q}$ is a pair $\langle m, i \rangle$, m is a cube over state variables, and $i \in \mathbb{N}$. A level N . A trace F_0, F_1, \dots .
Initially: $\mathcal{Q} = \emptyset, N = 0, F_0 = \text{Init}, \forall i > 0 \cdot F_i = \emptyset$.
repeat
 Unreachable If there is an $i < N$ s.t. $F_{i+1} \subseteq F_i$ **return** *Unreachable*.
 Reachable If there is an m s.t. $\langle m, 0 \rangle \in \mathcal{Q}$ **return** *Reachable*.
 Unfold If $F_N \rightarrow \neg \text{Bad}$, then set $N \leftarrow N + 1$.
 Candidate If for some $m, m \rightarrow F_N \wedge \text{Bad}$, then add $\langle m, N \rangle$ to \mathcal{Q} .
 Predecessor If $\langle m, i + 1 \rangle \in \mathcal{Q}$ and there are m_0 and m_1 s.t. $m_1 \rightarrow m, m_0 \wedge m_1'$ is satisfiable, and $m_0 \wedge m_1' \rightarrow F_i \wedge \text{Tr} \wedge m'$, then add $\langle m_0, i \rangle$ to \mathcal{Q} .
 NewLemma For $0 \leq i < N$: given a candidate model $\langle m, i + 1 \rangle \in \mathcal{Q}$ and clause φ , such that $\varphi \rightarrow \neg m$, if $\text{Init} \rightarrow \varphi$, and $\varphi \wedge F_i \wedge \text{Tr} \rightarrow \varphi'$, then add φ to F_j , for $j \leq i + 1$.
 ReQueue If $\langle m, i \rangle \in \mathcal{Q}, 0 < i < N$ and $F_{i-1} \wedge \text{Tr} \wedge m'$ is unsatisfiable, then add $\langle m, i + 1 \rangle$ to \mathcal{Q} .
 Push For $0 \leq i < N$ and a clause $(\varphi \vee \psi) \in F_i$, if $\varphi \notin F_{i+1}, \text{Init} \rightarrow \varphi$ and $\varphi \wedge F_i \wedge \text{Tr} \rightarrow \varphi'$, then add φ to F_j , for each $j \leq i + 1$.
until ∞ ;

Algorithm 1: IC3/PDR.

An *inductive trace*, or simply a trace, is a sequence of formulas $[F_0, \dots, F_k]$ that satisfy:

$$\text{Init} \Rightarrow F_0 \quad \forall 0 \leq i < k \cdot F_i(X) \wedge \text{Tr}(X, X') \rightarrow F_{i+1}(X') \quad (3)$$

An inductive trace is *monotone* if $\forall 0 \leq i < N \cdot F_i \Rightarrow F_{i+1}$ and *safe* when $\forall i \cdot F_i \Rightarrow \neg \text{Bad}$.

Note that for a monotone inductive trace $[F_0, \dots, F_k]$, F_i is a bounded invariant w.r.t. i for all $0 \leq i \leq k$.

3 IC3 and PDR

The finite state model checking algorithm IC3 was introduced in [?] and its variant PDR in [?]. It maintains sets of clauses $F_0, \dots, F_i, \dots, F_N$, called a *trace*, that are properties of states reachable in i steps from the initial states *Init*. Elements of F_i are called *lemmas*. In the following, we assume that F_0 is initialized to *Init*. After establishing that $\text{Init} \rightarrow \neg \text{Bad}$, the algorithm maintains the following invariants (for $0 \leq i < N$):

Invariant 1

$$F_i \rightarrow \neg \text{Bad} \quad F_i \rightarrow F_{i+1} \quad F_i \wedge \text{Tr} \rightarrow F_{i+1}'$$

That is, each F_i is safe, the trace is monotone, and F_{i+1} is inductive relative to F_i . In practice, the algorithm enforces monotonicity by maintaining $F_{i+1} \subseteq F_i$.

Alg. 1 summarizes, in a simplified form, a variant of the IC3 algorithm. The algorithm maintains a queue of counter-examples Q . Each element of Q is a tuple $\langle m, i \rangle$ where m is a monomial over v and $0 \leq i \leq N$. Intuitively, $\langle m, i \rangle$ means that a state m can reach a state in Bad in $N - i$ steps. Initially, Q is empty, $N = 0$ and $F_0 = Init$. Then, the rules are applied (possibly in a non-deterministic order) until either **Unreachable** or **Reachable** rule is applicable. **Unfold** rule extends the current trace and increases the level at which counterexample is searched. **Candidate** picks a set of bad states. **Predecessor** extends a counter-example from the queue by one step. **NewLemma** blocks a counterexample and adds a new lemma. **ReQueue** moves the counterexample to the next level. Finally, **Push** generalizes a lemma inductively. A typical schedule of the rules is to first apply all applicable rules except for **Push** and **Unfold**, followed by **Push** at all levels, then **Unfold**, and then repeating the cycle.

Queue. The queue is ordered by the level:

$$\langle m, i \rangle < \langle n, j \rangle \iff i < j \quad (4)$$

This drives the algorithm to the shortest counterexample.

Inductive Generalization. The **NewLemma** and **Push** rules are based on the principle of inductive generalization. Let $F_0, \dots, F_i, \dots, F_N$ be a valid trace, and let φ be a clause that is *relatively inductive* to F_i :

$$Init \implies \varphi \qquad \varphi \wedge F_i \wedge Tr \implies \varphi' \quad (5)$$

Let $G = G_0, \dots, G_N$ be defined as follows:

$$G_j = \begin{cases} F_j \cup \{\varphi\} & \text{if } j \leq i + 1 \\ F_j & \text{if } i + 1 < j \leq N \end{cases} \quad (6)$$

Then G is a valid trace. The proof is by induction on i and follows from monotonicity of the trace.

Generalizing predecessors. The **Predecessor** rule picks a predecessor m_0 in Tr of some (partial) state m . While it is possible to simply pick a predecessor state, the rule attempts to find a generalized predecessor instead. The conditions of the rule is sufficient to ensure that m_0 is an implicant of $\psi = (F_i \wedge \exists X' \cdot (Tr \wedge m'))$. Finding a prime implicant of ψ would have been even better, but is too expensive in practice.

Propagating lemmas. The **Push** rule propagates lemmas to higher level, optionally generalizing them as possible. This makes the trace “more” inductive, eventually leading to convergence.

Long counterexamples. The **ReQueue** rule lifts blocked counterexamples to higher levels. As a side-effect, it makes it possible to discover counterexamples longer than the current exploration bound N . For example, assume that m is blocked at level i . This means that there is a path of length $N - i$ from m to Bad (but no path of length at most i from $Init$ to m). Assume that **ReQueue** lifted m to level $j > i$, and then m was reachable from $Init$. Then, the discovered counterexample is a concatenation of a path of length k from $Init$ to m and a path of length $N - i$ from m to Bad . The total length of the counterexample is $(N - i + k)$ which is bigger than N .

4 *k*IC3

We present a variant of PDR which is designed for k -safety verification. For simplicity, we present the algorithm for $k = 2$. However, the extension to any k is natural.

Our goal is information flow analysis, we therefore first go over the definition of self-composition.

When using self-composition, information flow is tracked over an execution of two copies of the transition system, T and T_d . Let us denote $X_d := \{x_d \mid x \in X\}$ as the set of variables of T_d . Similarly, let $Init_d(X_d)$ and $Tr_d(X_d, X'_d)$ denote the initial states and transition relation of T_d . Given a formula φ over a set of variables V , $\varphi[V \leftarrow U]$ denotes the substitution of V with U in φ . Note that $Init_d$ and Tr_d are computed from $Init$ and Tr by means of substitutions. Namely, substituting every occurrence of $x \in X$ or $x' \in X'$ with $x_d \in X_d$ and $x'_d \in X'_d$, respectively. Formally, $Init_d := Init[X \leftarrow X_d]$ and $Tr_d[X \leftarrow X_d, X' \leftarrow X'_d]$. We formulate information flow over a self-composed program as a safety verification problem: $M_d := \langle Z, Init_d, Tr_d, Bad_d \rangle$ where $Z := X \cup X_d$ and

$$Init_d(Z) := Init(X) \wedge Init(X_d) \wedge \left(\bigwedge_{x \in L} x = x_d \right) \quad (7)$$

$$Tr_d(Z, Z') := Tr(X, X') \wedge Tr(X_d, X'_d) \quad (8)$$

$$Bad_d(Z) := \left(\bigvee_{x \in L} Obs_x(X) \wedge Obs_x(X_d) \wedge \neg(x = x_d) \right) \quad (9)$$

In order to track information flow, variables in L_d are initialized to be equal to their counterpart in L , while variables in H_d remain unconstrained. A leak is captured by the bad states (i.e. Bad_d). More precisely, there exists a leak iff there exists an execution of M_d that results in a state where $Obs_x(X)$, $Obs_x(X_d)$ hold and $x \neq x_d$ for a low-security variable $x \in L$.

The algorithm for k IC3 is summarized in Algorithm 2. Since this version aims at information flow analysis, the definition of Bad and $Init$ are incorporated implicitly into the algorithm as shown in the **Candidate**, **Predecessor**₀ and **Unfold** rules.

The main difference between IC3 and k IC3 is that the latter maintains a different queue, which tracks pairs (k tuple) of states.

Input: A transition system $\langle \text{Init}(X), \text{Tr}(X, X') \rangle$, a set of high vars $H \subseteq X$ and low variables $L = X \setminus H$.

Output: *Secure* or *Unsecure*

Data: A cex queue \mathcal{Q} , where $c \in \mathcal{Q}$ is a tuple $\langle m, n, i \rangle$, m and n are a cubes over state variables, and $i \in \mathbb{N}$. A level N . A trace F_0, F_1, \dots .

Initially: $\mathcal{Q} = \emptyset$, $N = 0$, $F_0 = \text{Init} \wedge \bigwedge_{x \in L} x = x_d, \forall i > 0 \cdot F_i = \emptyset$.

repeat

Unreachable If there is an $i < N$ s.t. $F_{i+1} \subseteq F_i$ **return** *Secure*.

Reachable If there is an m and n s.t. $\langle m, n, 0 \rangle \in \mathcal{Q}$ **return** *Unsecure*.

Unfold If $F_N(X) \wedge F_N(X_d) \rightarrow \bigwedge_{x \in L} x = x_d$, then set $N \leftarrow N + 1$.

Candidate If for some m and n , $m(X) \wedge n(X_d) \rightarrow F_N(X) \wedge F_N(X_d) \wedge \bigvee_{x \in L} x \neq x_d$, then add $\langle m, n[X_d \leftarrow X], N \rangle$ to \mathcal{Q} .

Predecessor If $\langle m, n, i + 1 \rangle \in \mathcal{Q}$ and there are (m_0, m_1) and (n_0, n_1) s.t. $m_1 \rightarrow m$ and $n_1 \rightarrow n$, $m_0 \wedge m'_1$ and $n_0 \wedge n'_1$ are satisfiable, and $m_0 \wedge m'_1 \rightarrow F_i \wedge \text{Tr} \wedge m'$ and $n_0 \wedge n'_1 \rightarrow F_i \wedge \text{Tr} \wedge n'$, then add $\langle m_0, n_0, i \rangle$ to \mathcal{Q} .

Predecessor₀ If $\langle m, n, 1 \rangle \in \mathcal{Q}$ and there are (m_0, m_1) and (n_0, n_1) s.t. $m_1 \rightarrow m$ and $n_1 \rightarrow n$, $m_0 \wedge m'_1$ and $n_0 \wedge n'_1$ are satisfiable, and $m_0 \wedge m'_1 \rightarrow F_0 \wedge \text{Tr} \wedge m'$ and $n_0 \wedge n'_1 \rightarrow F_0 \wedge \text{Tr} \wedge n'$, and $m_0|_L = n_0|_L$, then add $\langle m_0, n_0, 0 \rangle$ to \mathcal{Q} .

NewLemma For $0 \leq i < N$: given a candidate model $\langle m, n, i + 1 \rangle \in \mathcal{Q}$ and clause φ , such that $(\varphi \rightarrow \neg m) \vee (\varphi \rightarrow \neg n)$, if $\text{Init} \rightarrow \varphi$, and $\varphi \wedge F_i \wedge \text{Tr} \rightarrow \varphi'$, then add φ to F_j , for $j \leq i + 1$.

ReQueue If $\langle m, n, i \rangle \in \mathcal{Q}$, $0 < i < N$ and $F_{i-1} \wedge \text{Tr} \wedge (m \vee n)'$ is unsatisfiable, then add $\langle m, n, i + 1 \rangle$ to \mathcal{Q} .

Push For $0 \leq i < N$ and a clause $(\varphi \vee \psi) \in F_i$, if $\varphi \notin F_{i+1}$, $\text{Init} \rightarrow \varphi$ and $\varphi \wedge F_i \wedge \text{Tr} \rightarrow \varphi'$, then add φ to F_j , for each $j \leq i + 1$.

until ∞ ;

Algorithm 2: *kIC3*.